


ML Lab 3 Report

Question 1 :

1. Predictions using simple Decision Tree Regressor :

- Firstly, completed preprocessing, looked upon the dataset, there were no useless features, so did not drop any features, scaled the area values.
- And then built the Decision Regressor Tree code from scratch.



- Above is the graph for Y_original and Y_predicted comparison on testing data. We see for extreme high test price values, the model is failing to predict accurate prices.

Metrics for the DTR model on testing data :

MSE : 2126719553927.666016

RMS Error : 1458327.66

R2 score : 0.4227

2. 5-Fold Cross Validation :

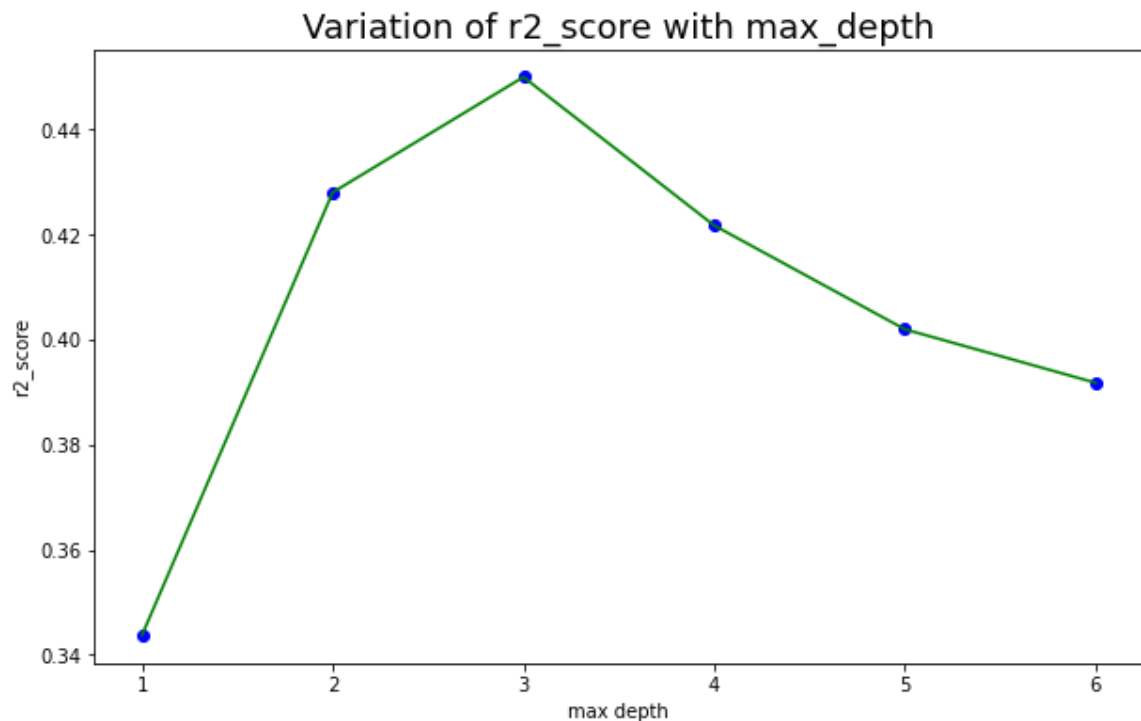
- Applied 5-fold cross validation on the whole training data and tested the variations with max_depth.
- Found optimal max_depth = 3, with maximum r2_score = 0.4499 .
- Tested the found optimal model on testing dataset and following were the results.

```
After 5-fold cross-validation and testing various max_depths,  
optimal max_depth = 3  
r2_score = 0.4499
```

```
Metrics of price predictions on Testing data :  
r2_score = 0.4438  
mean_absolute_error = 1012812.81
```

3. Visualisation of Cross Validation Results :





Hence we conclude that optimal max_depth = 3

4. Applying Bagging :

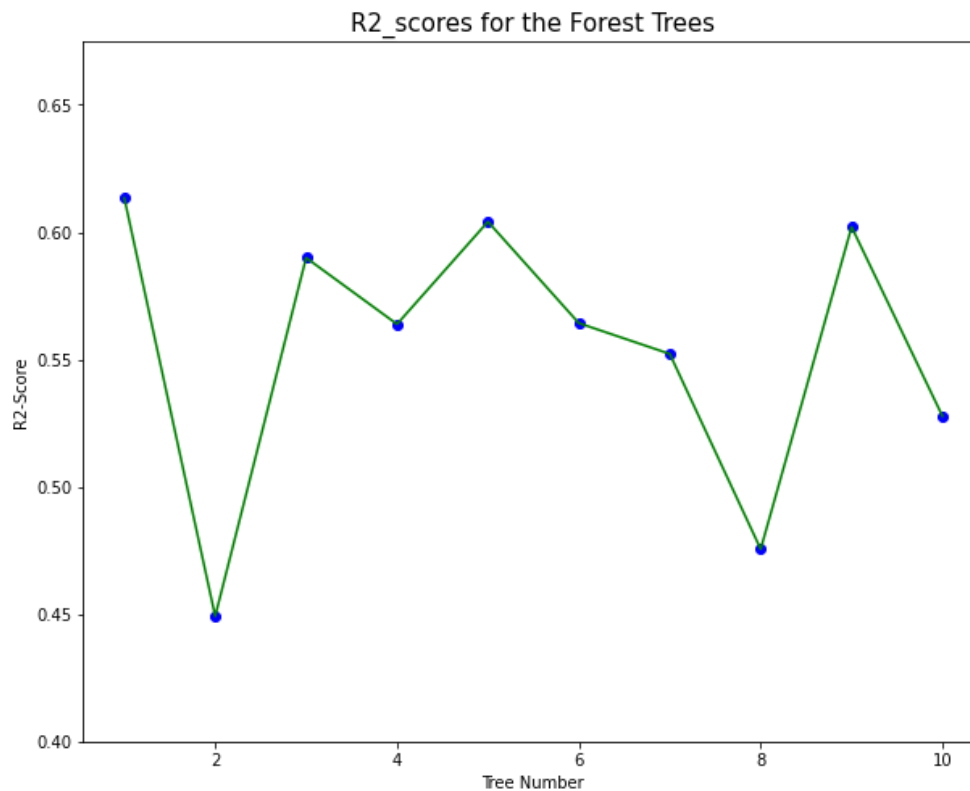
- For this task, we had to create a function which can in turn create different datasets by the BootStrap Aggregation method.

5. Training Decision Trees (Forest) :

- Now since we have created function that can create datasets, we can train Decision Trees on the created datasets.
- We have made a bagging() function to create 10 different models to train on the 10 datasets made by BootStrap Aggregation.

6. Performance of the Random Forest :

Average r2_score on BootStrapped data was 0.5542 . (varies when run again)



7. Combining the predictions of the Decision Trees :

- We have combined the predictions of the individual trees based on weights.
- The weights are scaled r2_scores of the respective trees.
- The predictions then made, and following are the results.

```
R-squared score on Testing Data, after combining the Random Forest Tree models :  
0.4761
```

We see that r2_score has improved,

For single DTR, r2_score = 0.4227

For Random Forest, r2_score = 0.4761 (varies when run again)

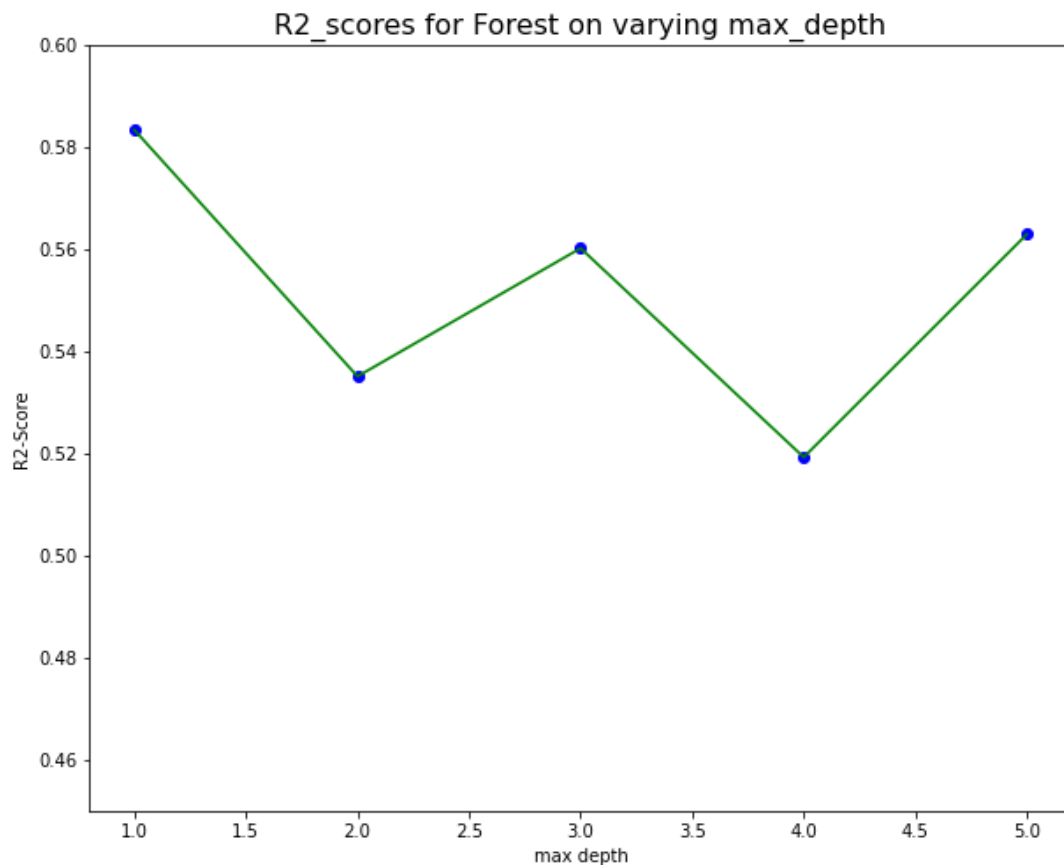
8. Varying max_depth :

Analysed the results when max_depth is varied on the forest model.

Here are the results

Average R-squared score for Combined Random Forest on Testing data with different max_depths :

```
max depth = 1 ; r2_score = 0.5833
max depth = 2 ; r2_score = 0.5351
max depth = 3 ; r2_score = 0.5602
max depth = 4 ; r2_score = 0.5192
max depth = 5 ; r2_score = 0.5629
```



I have run this many times, optimal max_depth varies on most of the times.

9. Training RandomForestRegressor from sklearn :

Following are the results for sklearn Random Forest.

For sklearn RandomForestRegressor, metrics on Testing data are

```
Mean Squared Error = 1964367302421.53
Mean Absolute Error = 988983.94
R-squared score = 0.4668
```

10. AdaBoost Regressor from sklearn :

Following are the results for sklearn AdaBoost.

```
For sklearn AdaBoostRegressor, metrics on Testing data are  
  
Mean Squared Error = 1709461030984.37  
Mean Absolute Error = 970131.22  
R-squared score = 0.5360
```

Question 2 :

1. Simple Decision Tree from Scratch :

Implemented Decision Tree Classifier from scratch, and evaluated it on Testing data. Following are the results.

```
Overall Accuracy for Trained Model : 92.40 %
```

2. 5-Fold Cross Validation on Simple DTC :

- Applied 5-fold cross validation on the simple DT classifier.
- Selected the best classifier model and evaluated it on the Testing dataset.
- Following were the results :

```
Progress Bar : 100%|██████████| 5/5 [00:41<00:00, 8.31s/it]  
After 5-fold cross-validation and testing various max_depths,  
optimal max_depth = 3  
Accuracy = 96.25 %  
  
Metrics of Classification on Testing data :  
Accuracy = 92.98
```

We see that Accuracy after cross validation and selecting the best model is 92.98 %. This is greater than previous model accuracy i.e. 92.40 %.

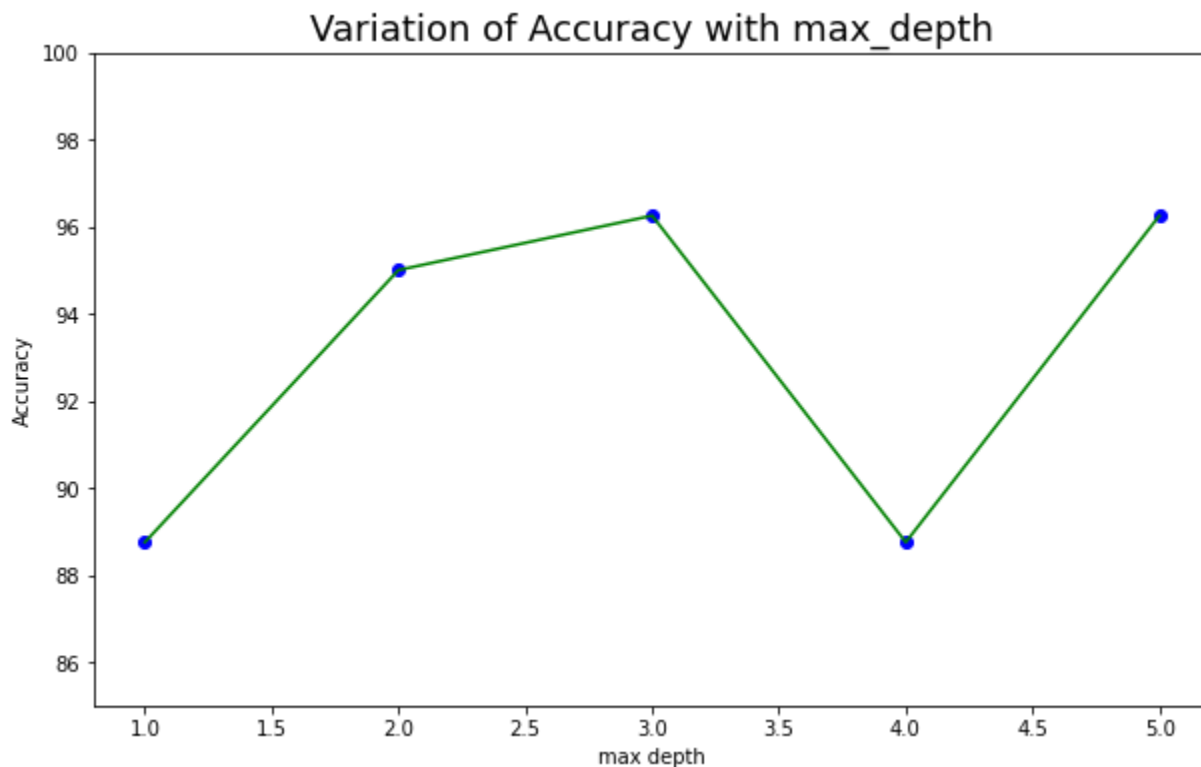
However after cross validation, we did not reach desired results, hence we go for boosting the Decision Trees with XGBoost and LightGBM.

3. Results of 5-fold Cross Validation :

Following are the results for cross validation on DTC :

```
Average Accuracies for Cross-Validation sets are :  
max_depth = 1 ; accuracy = 88.75  
max_depth = 2 ; accuracy = 95.00  
max_depth = 3 ; accuracy = 96.25  
max_depth = 4 ; accuracy = 88.75  
max_depth = 5 ; accuracy = 96.25
```

We find that max_depth = 3 gives maximum accuracy.



4. XGBoost Implementation :

Implemented XGBoost with subsample = 0.7, and max_depth = 4.

5. XGBoost Results :

We get the following results on the Testing Dataset :

```
Accuracy on Training Dataset : 98.74 %  
Accuracy on Testing Dataset : 95.32 %
```

6. LightGBM Implementation :

Implemented LightGBM with `max_depth = 3` and different `num_leaves`.


Following are results we get :

```
We get Maximum Accuracy on testing dataset when num_leaves = 3
```

```
Accuracies with num_leaves = 3 :  
Training Dataset Accuracy 95.98 %  
Testing Dataset Accuracy 95.32 %
```

7. Analysing the relation between max_depth and num_leaves :

We make following plot for analysing the relation :



- In the above charts, we can see that for **num_leaves ≤ 3** , training data classification accuracy, and testing data classification is increasing.
- But for **$4 \leq \text{num_leaves} < 8$** , testing accuracy keeps decreasing while training accuracy is still increasing. (initial overfitting)
- For **$8 \leq \text{num_leaves}$** , testing accuracy is minimum even though training accuracy is pretty good. (complete overfitting)

Hence we conclude that for **num_leaves $\geq 2^{(\text{max_depth})}$** , the model shows complete overfitting.

So we should usually choose **num_leaves $< 2^{(\text{max_depth})}$** .

8. Parameter Tuning :

We should control the leaf-wise tree growth for getting better accuracy on the testing dataset.

The most important parameters for tuning a LightGBM are :

1. `max_depth` :
We can control the max depth of the tree explicitly.
If we use a greater `max_depth`, we can overfit the tree, whereas if we keep it too low, we can underfit and have high bias. So optimum `max_depth` is chosen which is neither too low nor too high.
2. `num_leaves` :
We should choose a value which is sufficiently less than $2^{(\text{max_depth})}$ for `num_leaves`. For greater values, it will lead to overfitting as the tree grows complex.
3. `min_data_in_leaf` :
Its optimal value depends on the number of training samples and `num_leaves`. Setting it to a large value can avoid growing too deep a tree, but may cause under-fitting.

End of the Report !